


# Extreme Computing

## Hadoop

Stratis Viglas

School of Informatics  
University of Edinburgh  
`sviglas@inf.ed.ac.uk`

# Hadoop



Overview

Examples

Environment

# History

*Nutch* started in 2002 by Doug Cutting and Mike Cafarella

- ▶ Early open-source web-search engine
- ▶ Written in Java
- ▶ Realisation that it would not scale for the Web
  - ▶ 2004: Google MR and GFS papers appeared
  - ▶ 2005: Working MR implementation for Nutch
  - ▶ 2006: Hadoop became standalone project
- ▶ 2008: Hadoop broke world record for sorting 1TB of data

# Overview

Set of components (Java), implementing most of MR-related ecosystem:

- ▶ MapReduce
- ▶ HDFS (Hadoop distributed filesystem)
- ▶ Services on top:
  - ▶ HBase (BigTable)
  - ▶ Pig (sql-like job control)
- ▶ Job-control

# Overview

Hadoop supports a variety of ways to implement MR jobs:

- ▶ Natively, as java
- ▶ Using the 'streaming' interface
  - ▶ Mappers and reducers can be in any language
  - ▶ Performance penalty, restricted functionality
- ▶ C++ hooks etc

We will use the Streaming interface

## Aside: Map-Reduce in one line

Under Unix, you can quickly test a MR job:

```
% cat input | mapper | sort -0 +1 | reducer > output
```

**mapper** is your Mapper and **reducer** is the Reducer

# Word Counting

Word counting using Hadoop:

- ▶ Use HDFS
- ▶ Specify the MR program
- ▶ Run the job

Note: all commands are for Hadoop 0.19

# Word Counting

First need to upload data to HDFS

- ▶ Hadoop has a set of filesystem-like commands to do this:
  - ▶ **hadoop dfs -mkdir data**
  - ▶ **hadoop dfs -put file.text data/**
- ▶ This creates a new directory and uploads the file **file.txt** to HDFS
- ▶ We can verify that it is there:
  - ▶ **hadoop dfs -ls data/**



# Word Counting

## Mapper:

- ▶ Using Streaming, a Mapper reads from STDIN and writes to STDOUT
- ▶ Keys and Values are delimited (by default) using tabs.
- ▶ Records are split using newlines

# Word Counting

Mapper:

---

```
1 while !eof(STDIN) do  
2   | line = readLine(STDIN)  
3   | wordList = split(line)  
4   | foreach word in wordList do  
5   |   | print word TAB 1 NEWLINE  
6   | end  
7 end
```

---

# Word Counting

## Reducer

---

```
1 prevWord = "" ; valueTotal = 0
2 while !eof(STDIN) do
3     line = readLine(STDIN); (word,value) = split(line)
4     if word eq prevWord or prevWord eq "" then
5         valueTotal += value
6         prevWord = word
7     else
8         print prevWord valueTotal NEWLINE
9         prevWord = word; valueTotal = value
10    end
11 end
12 print word valueTotal NEWLINE
```

---

# Word Counting

## Improving the Mapper

---

```
1 wordsCounts = {}
2 while !eof(STDIN) do
3   | line = readLine(STDIN)
4   | wordList = split(line)
5   | foreach word in wordList do
6   |   | wordCounts[word]++
7   | end
8 end
9 foreach word in keys(wordCounts) do
10  | count = wordCounts[word]
11  | print word TAB count NEWLINE
12 end
```

---

# Word Counting

Our improved Mapper:

- ▶ Only emits one word-count pair, per word and shard
- ▶ This uses a *Combiner* technique
- ▶ Uses an unbounded amount of memory

Word counting 2 million tokens (Unix MR simulation)

Mapper	Time
Naive	1 minute 5 sec
Combiner	10 sec

How can you change it to use a *bounded* amount of memory?

# Secondary Sorting

At times, we may want to resort the Reducer input:

- ▶ Hadoop only guarantees that the same keys are grouped together
- ▶ We may want to ensure that some key occurs before other ones

*A secondary sort can rearrange the Reducer input*

## Secondary Sorting

Suppose we want to join two databases (White, p 234):

- ▶ Records in the DBs share the same set of IDs
- ▶ Want to merge records X in DB1 with records X in DB2
- ▶ All items in DB1 must be before DB2

## Secondary Sorting

DB1		DB2		
Station ID	Station Name	Station ID	Time	Temperature
A1	Zebra	A1	10:00	12
A2	Goat	A1	11:00	11
		A1	12:00	9
		A2	10:00	12

Combined:

Station ID	Time	Station Name	Temperature
A1	10:00	Zebra	12
A1	11:00	Zebra	11
A1	12:00	Zebra	9
A2	10:00	Goat	12



# Secondary Sorting

Approach:

- ▶ Key mapper output by record ID
- ▶ Assign a secondary key to each output
- ▶ Resort output by first and second key
  - ▶ All records with the same ID are still together
  - ▶ Now, DB1 can be before DB2

# Running a job

To run a streaming job:

```
hadoop jar \  
contrib/streaming/hadoop-*-streaming.jar \  
-mapper map -reducer red -input in \  
-output out -file map -file red
```

- ▶ **map** and **red** are the actual MR program files
- ▶ **in** is the input DFS and **out** is the output DFS directory

# Job Control

There is a graphical UI:

**<http://crom.inf.ed.ac.uk:50030/jobtracker.jsp>**

And also a command-line interface:

- ▶ **`hadoop job -list`**
- ▶ **`hadoop job -status jobid`**
- ▶ **`hadoop job -kill jobid`**

# Summary

- ▶ History
- ▶ Looked at major components
- ▶ Two examples
- ▶ Job control