

Extreme Computing

Map-Reduce 2

Stratis Viglas

School of Informatics
University of Edinburgh
sviglas@inf.ed.ac.uk

Programming Model

Examples

Efficiency

Programming Model

MR offers one restricted version of parallel programming:

- ▶ Coarse-grained.
- ▶ No inter-process communication.
- ▶ Communication is (generally) through files.

Programming Model

Mapping:

- ▶ The input data is divided into *shards*.
- ▶ The *Map* operation works over each shard and *emits key-value* pairs.
- ▶ Each mapper works in parallel.


Keys and values can be anything which can be represented as a string.

Programming Model

Reducing:

- ▶ After mapping, each key-value pair is hashed on the key.
- ▶ Hashing sends that key-value pair to a given *reducer*.
 - ▶ All keys that hash to the same value are sent to the same reducer.
- ▶ The input to a reducer is sorted on the key.
 - ▶ Sorted input means that related key-value pairs are locally grouped together.

Programming Model



Programming Model

Note:

- ▶ Each Mapper and Reducer runs in parallel.
- ▶ There is no state sharing between tasks.
 - ▶ Task communication is achieved using either external resources or at start-time
- ▶ There need not be the same number of Mappers as Reducers.
 - ▶ It is possible to have *no* Reducers.

Programming Model

Note:

- ▶ Tasks read their input sequentially.
 - ▶ Sequential disk reading is far more efficient than random access
- ▶ Reducing starts once Mapping ends.
 - ▶ Sorting and merging etc can be interleaved.

Example: System Maintenance

Example

Check the health of a set of machines. Report on any that are bad.

- ▶ The input data will be a list of machine names
 - ▶ borg1.com, borg2.com, ... borg12123.com
- ▶ The output will be reports
- ▶ There is no need to run a reducing stage

Example: System Maintenance

Mapper:

- ▶ Read each line, giving a machine name,
- ▶ Log into that machine and collect statistics.
- ▶ Emit statistics

Example: Word Counting

Example

Count the number of words in a collection of documents

- ▶ Our Mapper counts words in each shard.
- ▶ The Reducer gathers together partial counts for a given word and sums them

Example: Word Counting

Mapper:

- ▶ For each sentence, emit *word*, *1* pair.
 - ▶ The key is the *word*
 - ▶ The value is the number 1

Example: Word Counting

Reducer:

- ▶ Each Reducer will see all instances of a given word.
- ▶ Sequential reads of the reducer input give partial counts of a word.
- ▶ Partial counts can be summed to give the total count.

Example: Word Counting

Input sentences:

- ▶ *the cat*
- ▶ *the dog*

Key	Value	
<i>the</i>	1	
<i>cat</i>	1	Mapper output
<i>the</i>	1	
<i>dog</i>	1	

Example: Word Counting

Reducer 1 input

the, 1

the, 1

dog, 1

Reducer 2 input

cat, 1

Reducer 1 output

the, 2

dog, 1

Reducer 2 output

cat, 1

Map Reduce Efficiency

MR algorithms involve a lot of disk and network traffic:

- ▶ We typically start with Big Data
- ▶ Mappers can produce intermediate results that are *bigger* than the input data.
- ▶ Task input may not be on the same machine as that task.
 - ▶ This implies network traffic
- ▶ Per-reducer input needs to be sorted.

Map Reduce Efficiency

Sharding might not produce a balanced set of inputs for each Reducer:

- ▶ Often, the data is heavily skewed
 - ▶ Eg all function words might go to one Reducer
- ▶ Having an imbalanced set of inputs turns a parallel algorithm into a sequential one

Map Reduce Efficiency

Selecting the right number of Mappers and Reducers can improve speed

- ▶ More tasks mean each task might fit in memory / require less network access
- ▶ More tasks mean that failures are quicker to recover from.
- ▶ Fewer tasks have less of an over-head.

This is a matter of guess-work

Map Reduce Efficiency

Algorithmically, we can:

- ▶ Emit fewer key-value pairs
 - ▶ Each task can locally aggregate results and periodically emit them.
 - ▶ (This is called *combining*)
- ▶ Change the key
 - ▶ Key selection implies we partition the output. Some other selection might partition it more evenly

Summary

- ▶ Introduced the MR programming model
- ▶ Sample MR applications
- ▶ Looked at efficiency